



Do phonology-derived ventral representations occupy the same semantic space as their GloVe targets?

One 2-component PCA fitted on the L2-normalised GloVe bank ONLY, then applied unchanged to GloVe and to each seed's $\hat{s}$ (PC1 + PC2 = 8.3 % of GloVe variance)

seed 19 / epoch 155



seed 22 / epoch 140



● target GloVe, all 29,571 words ● predicted $\hat{s}$ (L2-normalised), all 29,571 ● annotated word: GloVe target ■ annotated word: predicted $\hat{s}$ — target -> prediction